

# Newt, assemble web applications with Pandoc, Postgres and PostgREST

R. S. Doiel, [rsdoiel@caltech.edu](mailto:rsdoiel@caltech.edu)

Caltech Library, Digital Library Development

July 14, 2023



# Today, LAMP and its legacy

Four example systems found in Caltech Library

- ▶ EPrints
- ▶ Archivesspace
- ▶ Islandora
- ▶ Invenio RDM

## Required Knowledge

App	Languages	Supporting services
ArchivesSpace	Java, Ruby, SQL	MySQL, Solr, Apache or NginX, Solr
EPrints	Perl, SQL, XML, EPrint templates	MySQL, Apache2 (tight integration), Sphynx
Invenio RDM	Python,SQL JavaScript/React NodeJS/NPM	Postgres, Redis, Elasticsearch, Docker, Invenio Framework, Python packaging system
Islandora	PHP/SQL	MySQL, Fedora, Apache 2

These are all really complicated pieces of software.

## The problem

Each listed application is built on a stack. The stacks are complex. Because of the complexity it's hard to sustain them. Some we've outsourced to SAAS providers (e.g. ArchivesSpace). Some we treat as a back boxes (e.g. EPrints). It's just not fun supporting applications at this level of complexity. It takes too much time and energy. It detracts from delivering useful things to our Library, Archives and Caltech Community.

# Why are these things so complex?

*WARNING: gross generalizations ahead*

- ▶ We want more from our application so more code gets written, complexity accrues over time
- ▶ We must build systems to scale!
  - ▶ a not so subtle influence on developer “best practices” from Silicon Valley

## Let's talk about scale

- ▶ Best practice often translates to building for scale, specifically scaling up
- ▶ Scaling up => programmable infrastructure, the siren song of Google, AWS and Azure
  - ▶ Scaling big is hard
  - ▶ Scaling big makes things really complex
  - ▶ Do we ever really need to build at Google/Amazon/Azure scale?

# The alternative

- ▶ Scaling down  $\leq$  pack only what you need
- ▶ Simplify!



# Scaling down

- ▶ Limit the moving parts
- ▶ Limit the cognitive shifts
- ▶ Minimize the toolbox while maximizing how you use it
- ▶ Write less code! But remain readable!

# How minimal can we go?

- ▶ Off the self microservices
- ▶ Build with SQL and Pandoc

# Can you make web applications using only SQL and Pandoc?

Just about. Here's the off the shelf microservices I am experimenting with

- ▶ Postgres
- ▶ PostgREST
- ▶ Pandoc
- ▶ Newt

# A clear division of labor

Simplify through a clear division of labor.

- ▶ `Postgres` + `PostgREST` => JSON DATA API
- ▶ `Pandoc` => Powerful template engine
- ▶ `Newt` => data router

## How would this work in practice?

1. Model our data using SQL (Postgres)
2. Define our JSON API using SQL (Postgres+PostgREST)
3. Transform our structured data using Pandoc
4. Use Newt to orchestrate

A game of telephone web browser => Newt => PostgREST => Pandoc =>  
Newt => web browser

# Required Toolbox

- ▶ Text editor
- ▶ Spreadsheet (optional)
- ▶ Web browser
- ▶ Pandoc
- ▶ Postgres + PostgREST
- ▶ Newt

# Server side knowledge requirements

- ▶ SQL
- ▶ Pandoc templates
- ▶ CSV file describing data flowing through the microservices

# Web browser knowledge requirements

- ▶ HTML
- ▶ CSS (optional)
- ▶ JavaScript (optional)



# What does this enable?

We can create interactive applications with some SQL, Pandoc templates and a little routing.

## Why SQL?

SQL is really good at describing structured data. It also is good at expressing queries. With a little deeper knowledge of SQL you can also define data views, functions and your own procedures. With Postgres + PostgREST these provide everything you need in a JSON data API short of file upload. SQL does a lot of lifting with a little code and usually remains readable.

*Minimize the source Luke!*

You don't need to design classes in your favorite programming languages and Schema in SQL. You don't need to learn an ORM. You don't duplicate the code in the data base, again in the middleware and increasingly often in the browser. Data models are defined in one place, Postgres. PostgREST takes care of turning them into a JSON data API. Data transformation is handled by Pandoc. A program really good at translating things. Newt provides just enough orchestration based on defining some routes in a CSV file.

## Fewer cognitive shifts

- ▶ SQL (Postgres dialect)
- ▶ JSON
- ▶ Pandoc templates
- ▶ Describing data flow in a CSV file using a simple notation

data flow web browser => Newt => PostgREST => Pandoc => Newt =>  
web browser

## helpful to know

- ▶ HTML 5 related W3C technologies
  - ▶ HTML 5 markup
  - ▶ CSS
  - ▶ JavaScript
- ▶ Understand how HTTP works, including HTTP methods and Headers
- ▶ How to handle static file assets, e.g. image and video files

But is this really simpler?

# What was needed for each version of birds?

Let's take a look at three versions of a bird sighting website.

- ▶ **birds 1**, a static site implementation
- ▶ **birds 2**, a dynamic site implementation, content viewing requires browser JavaScript
- ▶ **birds 3**, a dynamic site implementation, does not require browser JavaScript

# Different birds

## birds 1 static site (read only)

- ▶ Built with Pandoc from Markdown and CSV file
- ▶ Adds bird sightings via updating a CSV file and rebuilding site with Pandoc

# Different birds

## birds 2, dynamic site (read/write)

- ▶ Built with SQL using Postgres + PostgREST
- ▶ Add birds you using a web form
- ▶ Requires the web browser to assemble pages via API calls
- ▶ JavaScript becomes complex between fetching data and inserting it into the page
- ▶ Doesn't work in text only web browsers like Lynx



# Different birds

## birds 3, dynamic site (read/write)

- ▶ Build from SQL (Postgres + PostgREST) and Pandoc
- ▶ Add birds you using a web form
- ▶ Rendered on server and no longer requires JavaScript
- ▶ Works even in text web browsers like Lynx

# Different birds

## Pros and cons

version	site type	pros/cons
birds 1	static	easy to conceptualize / read only
birds 2	dynamic	read write site / requires browser JavaScript, JavaScript is complex
birds 3	dynamic	read write site, easy to conceptualize / requires SQL and knowledge of Pandoc

## Birds 3 => Postgres+PostgREST, Pandoc and Newt

*The complicated activities are handled by the off the self microservices.  
The remaining complexity is limited to SQL to model data and our Pandoc templates.*

- ▶ Avoids browser side page assembly
- ▶ Leverages our Pandoc knowledge
- ▶ Data is modeled using SQL

## Newt manages data flow

- ▶ request => data API => Pandoc => response
- ▶ Newt's routes can be managed in spreadsheet!

# Developer workflow

1. Model data in Postgres
2. Create/edit Pandoc templates
3. Create/edit routes CSV file in a spreadsheet
4. (Re)start Newt and PostgREST to (re)load models and routes

**Repeat as needed**

# Minimizing newness

- ▶ If you've attended a data science workshop you are likely know enough SQL
- ▶ If you've built a static website with Pandoc you know how Pandoc works
- ▶ I think there is community that knows some SQL, CSV files and knows Pandoc

=> Is this useful for Libraries, Archives and Museums?

# Evaluating Postgres+PostgREST, Pandoc and Newt

## ▶ Weaknesses

- ▶ Newt is limited (doesn't support file uploads)
- ▶ PostgREST is new to me (and probably others)
- ▶ SQL and HTML have a learning curve

## ▶ Strengths

- ▶ SQL is proven and likely to be around a very long time
- ▶ HTML is proven and likely to be around a very long time
- ▶ Postgres and Pandoc are very mature software
- ▶ PostgREST and Newt are useful today and worth exploring

## Next steps for Newt?

- ▶ Newt is an experiment, I am building some staff applications Summer/Fall 2023
- ▶ Solr/Opensearch should work with Pandoc and Newt, seems promising

somday, maybe

- ▶ It'd be nice if Newt could send file uploads to a service like S3 (via Minio?)
- ▶ It might be nice if Newt could function as a static file server



# Thank you!

- ▶ Presentation <https://caltechlibrary.github.io/newt/presentation/>
- ▶ Project: <https://github.com/caltechlibrary/newt>
- ▶ Email: [rsdoiel@caltech.edu](mailto:rsdoiel@caltech.edu)